

# QEMU'ing up a storm

Why QEMU is pretty goddamn awesome!



# whoami

- ▶ Nominally do payment security stuff - i.e PCI...
- ▶ Sometimes do silly stuff
- ▶ Email: [peter@peterfillmore.com](mailto:peter@peterfillmore.com)
- ▶ Github: <http://www.github.com/peterfillmore>

# How I got to QEMU









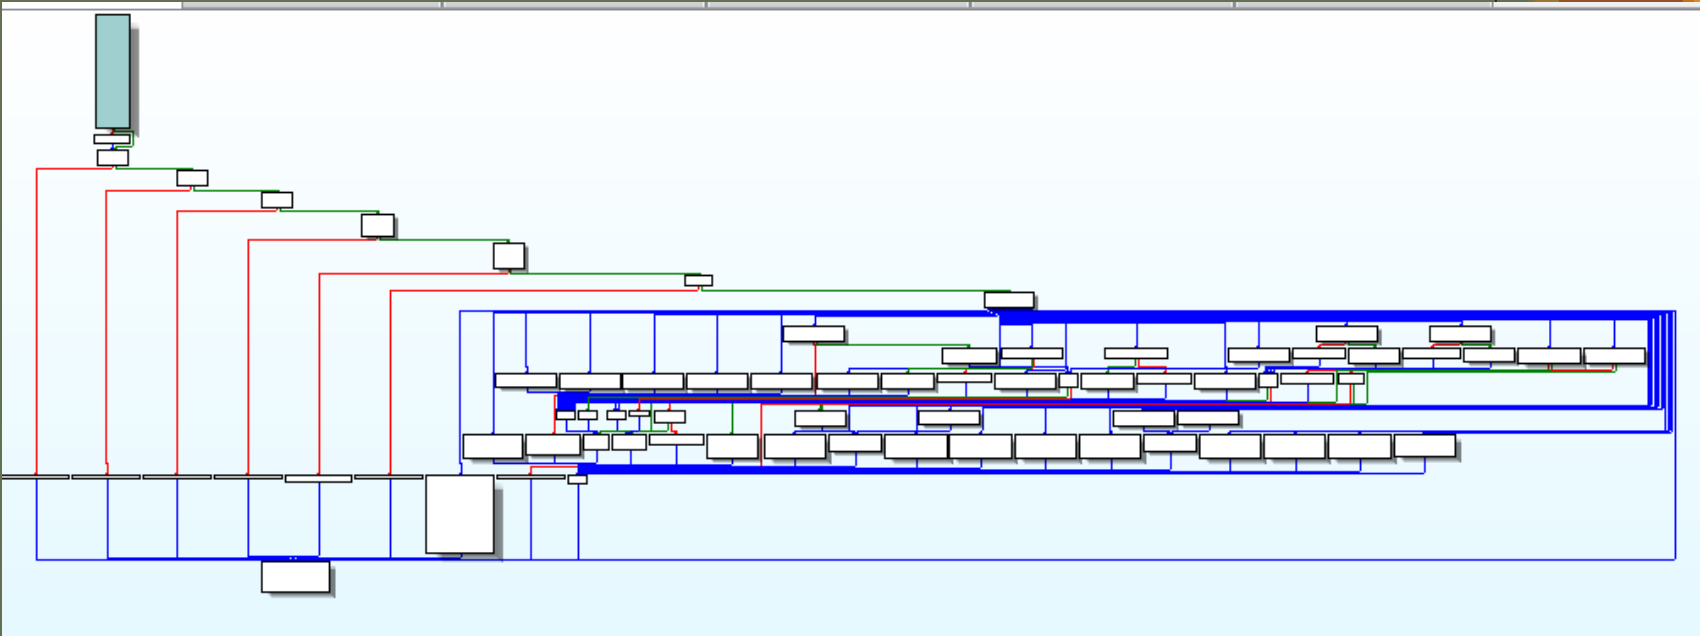
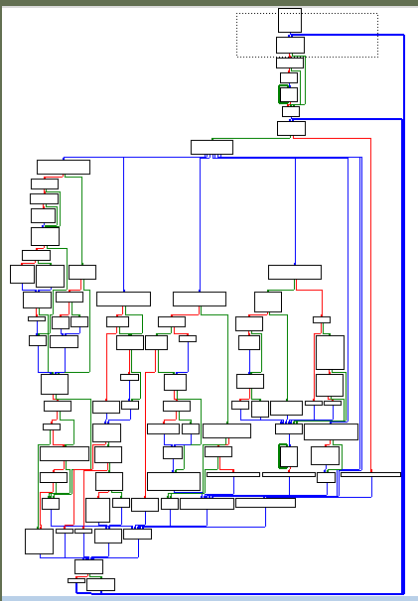






00000000								18 F0 9F E5	18 F0 9F E5	18 F0 9F E5	18 F0 9F E5	
00000008	18 F0 9F E5	18 F0 9F E5	18 F0 9F E5	18 F0 9F E5	18 F0 9F E5	18 F0 9F E5	18 F0 9F E5	80 5F 20 B9	.=■s.=■s.=■sÇ_ !			
00000018	F0 FF 1F E5	18 F0 9F E5	18 F0 9F E5	18 F0 9F E5	18 F0 9F E5	18 F0 9F E5	18 F0 9F E5	58 00 00 00	= .s.=■sX...@...			
00000028	44 00 00 00	48 00 00 00	48 00 00 00	48 00 00 00	48 00 00 00	48 00 00 00	48 00 00 00	4C 00 00 00	D...H...L.....			
00000038	50 00 00 00	54 00 00 00	54 00 00 00	54 00 00 00	54 00 00 00	54 00 00 00	54 00 00 00	FE FF FF EA	P...T...  0  0			
00000048	FE FF FF EA	FE FF FF EA	FE FF FF EA	FE FF FF EA	FE FF FF EA	FE FF FF EA	FE FF FF EA	FE FF FF EA	0  0  0  0			
00000058	DC 00 9F E5	AA 10 A0 E3	AA 10 A0 E3	AA 10 A0 E3	AA 10 A0 E3	AA 10 A0 E3	AA 10 A0 E3	55 20 A0 E3	_■s¬.ápU áp%0áp			
00000068	04 30 80 E5	01 30 A0 E3	01 30 A0 E3	01 30 A0 E3	01 30 A0 E3	01 30 A0 E3	01 30 A0 E3	00 30 80 E5	.0Çs.0áp.0Çs..Çs			
00000078	0C 20 80 E5	08 30 90 E5	08 30 90 E5	08 30 90 E5	08 30 90 E5	08 30 90 E5	08 30 90 E5	01 38 13 E2	. Çs.0És.;.Gn .			
00000088	03 30 A0 E3	00 30 80 E5	00 30 80 E5	00 30 80 E5	00 30 80 E5	00 30 80 E5	00 30 80 E5	0C 10 80 E5	.0áp.0Çs..Çs. Çs			
00000098	A0 00 9F E5	04 10 A0 E3	04 10 A0 E3	04 10 A0 E3	04 10 A0 E3	04 10 A0 E3	04 10 A0 E3	04 10 80 E5	á.■s..áp..Çs..áp			
000000A8	00 10 80 E5	90 00 9F E5	90 00 9F E5	90 00 9F E5	90 00 9F E5	90 00 9F E5	90 00 9F E5	DB F0 21 E3	..ÇsÉ.■s =!p.-áb			
000000B8	04 00 40 E2	D7 F0 21 E3	D7 F0 21 E3	D7 F0 21 E3	D7 F0 21 E3	D7 F0 21 E3	D7 F0 21 E3	00 D0 A0 E1	..@G+=!p.-áb..@G			
000000C8	D1 F0 21 E3	00 D0 A0 E1	00 D0 A0 E1	00 D0 A0 E1	00 D0 A0 E1	00 D0 A0 E1	00 D0 A0 E1	04 00 40 E2	-=!p.-áb..@G-=!p			
000000D8	00 D0 A0 E1	01 0B 40 E2	01 0B 40 E2	01 0B 40 E2	01 0B 40 E2	01 0B 40 E2	01 0B 40 E2	D3 F0 21 E3	.-áb..@G+=!p.-áb			
000000E8	04 00 40 E2	10 F0 21 E3	10 F0 21 E3	10 F0 21 E3	10 F0 21 E3	10 F0 21 E3	10 F0 21 E3	00 D0 A0 E1	..@G.=!p.-áb.½MG			
000000F8	48 10 9F E5	48 20 9F E5	48 20 9F E5	48 20 9F E5	48 20 9F E5	48 20 9F E5	48 20 9F E5	48 30 9F E5	H.■sH ■sH0■s..R0			
00000108	04 00 91 34	04 00 82 34	04 00 82 34	04 00 82 34	04 00 82 34	04 00 82 34	04 00 82 34	FB FF FF 3A	..æ4..é4v :..áp			
00000118	34 10 9F E5	34 20 9F E5	34 20 9F E5	34 20 9F E5	34 20 9F E5	34 20 9F E5	34 20 9F E5	02 00 51 E1	4.■s4 ■s..Q0..ü4			
00000128	FC FF FF 3A	04 E0 8F E2	04 E0 8F E2	04 E0 8F E2	04 E0 8F E2	04 E0 8F E2	04 E0 8F E2	24 00 9F E5	n :.aǺG\$.■s. /0			
00000138	FE FF FF EA	80 C0 1F E0	80 C0 1F E0	80 C0 1F E0	80 C0 1F E0	80 C0 1F E0	80 C0 1F E0	00 C0 1F E0	0Ç+.a.+ .a.Ç.@			
00000148	80 29 05 00	00 00 00 40	00 00 00 40	00 00 00 40	00 00 00 40	00 00 00 40	00 00 00 40	FC 0B 00 40	Ç).....@n..@...@			





# QEMU

- ▶ “Quick Emulator”
- ▶ Originally written by Fabrice Bellard
- ▶ That dude is scary good, you’ve heard of:
  - ▶ FFMPEG
  - ▶ 4G LTE Base Station  
<http://bellard.org/lte/>
  - ▶ LZEXE !?!  
<http://bellard.org/lzexe.html>  
<http://bellard.org/>

# QEMU - Why so fast?

- ▶ Dynamic translation of target instructions to host.
- ▶ Completely written in C
- ▶ Uses TCG - Tiny Code Generator to generate RISC like instructions.
- ▶ These are then compiled dynamically for the host.
- ▶ So, speeeeeeed.



# What can you use it for?

- ▶ User Mode Emulation - emulate a program compiled in another language directly on the host.
- ▶ System Emulation - emulate a complete system
- ▶ Kernel-based Virtual Machine stuff
- ▶ Xen Hosting stuff

# What can I emulate?

PC (x86 or x86_64 processor)	Luminary Micro LM3S811EVB (ARM Cortex-M3)
PREP (PowerPC processor)	Luminary Micro LM3S6965EVB (ARM Cortex-M3)
ISA PC (old style PC without PCI bus)	Freescale MCF5208EVB (ColdFire V2).
G3 Beige PowerMac (PowerPC processor)	Arnewsh MCF5206 evaluation board (ColdFire V2).
Mac99 PowerMac (PowerPC processor, in progress)	Palm Tungsten E PDA (OMAP310 processor)
Sun4m/Sun4c/Sun4d (32-bit Sparc processor)	N800 and N810 tablets (OMAP2420 processor)
Sun4u/Sun4v (64-bit Sparc processor, in progress)	MusicPal (MV88W8618 ARM processor)
Malta board (32-bit and 64-bit MIPS processors)	Gumstix "Connex" and "Verdex" motherboards (PXA255/270).
MIPS Magnum (64-bit MIPS processor)	Siemens SX1 smartphone (OMAP310 processor)
ARM Integrator/CP (ARM)	AXIS-Devboard88 (CRISv32 ETRAX-FS).
ARM Versatile baseboard (ARM)	Spitz, Akita, Borzoi, Terrier and Tosa PDAs (PXA270 processor)
ARM RealView Emulation/Platform baseboard (ARM)	Petalogix Spartan 3aDSP1800 MMU ref design (MicroBlaze).

# Why can't I just QEMU out of the box?

- ▶ QEMU is set out of the box for running OS's
- ▶ However my image has no OS - it's bare metal
- ▶ I know I have an ARM - so copy the "versatilePB.c" to "vivotech.c" - and add to make files
- ▶ Read QEMU code - get jealous as its pretty good - "self documenting"
- ▶ Make some stupid errors!



4.0 GB	AHB PERIPHERALS	0xFFFF FFFF
3.75 GB		0xF000 0000
	APB PERIPHERALS	
3.5 GB		0xE000 0000
	RESERVED ADDRESS SPACE	
3.0 GB		0xC000 0000
2.0 GB	BOOT BLOCK (RE-MAPPED FROM ON-CHIP FLASH MEMORY)	0x8000 0000
	RESERVED ADDRESS SPACE	
		0x4001 8000
		0x4000 7FFF
	TOTAL OF 32 kB ON-CHIP STATIC RAM (LPC2136/38)	0x4000 4000
	TOTAL OF 16 kB ON-CHIP STATIC RAM (LPC2132/34)	0x4000 3FFF
		0x4000 2000
	TOTAL OF 8 kB ON-CHIP STATIC RAM (LPC2131)	0x4000 1FFF
1.0 GB		0x4000 0000
	RESERVED ADDRESS SPACE	
		0x0008 0000
	TOTAL OF 512 kB ON-CHIP NON-VOLATILE MEMORY (LPC2138)	0x0007 FFFF
		0x0004 0000
	TOTAL OF 256 kB ON-CHIP NON-VOLATILE MEMORY (LPC2136)	0x0003 FFFF
		0x0002 0000
	TOTAL OF 128 kB ON-CHIP NON-VOLATILE MEMORY (LPC2134)	0x0001 FFFF
		0x0001 0000
	TOTAL OF 64 kB ON-CHIP NON-VOLATILE MEMORY (LPC2132)	0x0000 FFFF
		0x0000 8000
	TOTAL OF 32 kB ON-CHIP NON-VOLATILE MEMORY (LPC2131)	0x0000 7FFF
0.0 GB		0x0000 0000

# Initializing your RAM

```
28 #define VIVOTECH_RAM_ADDR 0x40000000
29 #define VIVOTECH_RAM_SIZE (64 * 1024)
```

```
111     memory_region_init_ram(ram, NULL, "vivotech.ram", machine->ram_size,
112                             &error_abort);
113     vmstate_register_ram_global(ram);
114     /* ??? RAM should repeat to fill physical memory space. */
115     /* SDRAM at address zero. */
116     memory_region_add_subregion(sysmem, VIVOTECH_RAM_ADDR, ram);
```

# Initializing your ROM

```
25 #define VIVOTECH_FLASH_ADDR 0x00000000
26 #define VIVOTECH_FLASH_SIZE (512 * 1024)
27 #define VIVOTECH_FLASH_SECT_SIZE 512
```

```
247     dinfo = drive_get(IF_PFLASH, 0, 0);
248     if (!pflash_cfi01_register(VIVOTECH_FLASH_ADDR, NULL, "vivotech.flash",
249                               VIVOTECH_FLASH_SIZE,
250                               dinfo ? blk_by_legacy_dinfo(dinfo) : NULL,
251                               VIVOTECH_FLASH_SECT_SIZE,
252                               VIVOTECH_FLASH_SIZE / VIVOTECH_FLASH_SECT_SIZE,
253                               4, 0x0089, 0x0018, 0x0000, 0x0, 0)) {
254         fprintf(stderr, "qemu: Error registering flash memory.\n");
255     }
256
```



# Shell Code for ROM functions

```
//shim for IAP function on the LPC213x

void iapfunction()
{
    asm("ldr r1, [r0]");
    asm("cmp r1, #54");
    asm("ldr r2, [pc, #12]");
    asm("str r2, [r4, #4]");
    register unsigned long *commandparameter asm("r0");
    register unsigned long *commandresult asm("r4");
    commandresult[0] = 0x00; //always return 0
    if(commandparameter[0] == 0x36){
        commandresult[1] = 0x2FF25; //LPC2138 part id
    }
}
```

```
37 char iapcode[41] = {
38 0x01, 0x68, 0x36, 0x29,
39 0xc0, 0x46, 0xc0, 0x46,
40 0xc0, 0x46, 0xc0, 0x46,
41 0xc0, 0x46, 0xc0, 0x46,
42 0x03, 0xd1, 0x03, 0x4a,
43 0x62, 0x60, 0x70, 0x47,
44 0xc0, 0x46, 0x70, 0x47,
45 0xc0, 0x46, 0xc0, 0x46,
46 0x25, 0xff, 0x02, 0x00};
```

```
142 memory_region_init_ram(iap, NULL, "vivotech.iap", 0x2000,
143                          &error_abort);
144 vmstate_register_ram_global(iap);
145 memory_region_add_subregion(sysmem, 0x7FFFF000, iap);
146 //copy data into the memory address
147 char *ramptr = (char *)memory_region_get_ram_ptr(iap);
148 if(ramptr != NULL)
149     memcpy(ramptr+0xFF0, iapcode, sizeof(iapcode));
```

# QEMU Command Line

```
1 qemu-system-arm -M vivotech -pflash flash.img -m 32k -nographic -S -s
```

- ▶ -M = machine - in this case its our vivotech platform
- ▶ -pflash = our binary flash image
- ▶ -m = guest ram (in our case 32k)
- ▶ -nographic = no display
- ▶ -S = freeze CPU at startup
- ▶ -s = start the GDB server

# So why not just use hardware?

- ▶ Full control over execution - can single step, read memory, insert code anywhere.
- ▶ Hardware may not have debugging enabled.
- ▶ Hardware may be slow and resource constrained
- ▶ Can run multiple instances on the same machine

# Future Stuff I want to do

- ▶ Convince SoC vendors to ship their own QEMU image platforms for testing.
- ▶ Use QEMU to set up fuzzing farms for embedded systems
- ▶ Slowly working on implementing basic crash handler for embedded arm systems